← Back **Continuous state spaces**
Graded Assignment • 30 min

Due Apr 8, 11:59 PM CEST

Your grade: 100%

Your latest: **100%** • Your highest: **100%** • To pass you need at least 80%. We keep your highest score.

Next item →

1 / 1 point

1. The Lunar Lander is a continuous state Markov Decision Process (MDP) because:

- (●) The state contains numbers such as position and velocity that are continuous valued.
- ( ) The state-action value $Q(s, a)$ function outputs continuous valued numbers
- ( ) The reward contains numbers that are continuous valued
- ( ) The state has multiple numbers rather than only a single number (such as position in the x-direction)

Correct
That's right!

1 / 1 point

2. In the learning algorithm described in the videos, we repeatedly create an artificial training set to which we apply supervised learning where the input $x = (s, a)$ and the target, constructed using Bellman's equations, is y = _____?

- ( ) $y = R(s')$ where s' is the state you get to after taking action a in state s
- (●) $y = R(s) + \gamma \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s
- ( ) $y = R(s)$
- ( ) $y = \max_{a'} Q(s', a')$ where s' is the state you get to after taking action a in state s

Correct

1 / 1 point

3. You have reached the final practice quiz of this class! What does that mean? (Please check all the answers, because all of them are correct!)

- [x] What an accomplishment -- you made it!

  Correct

- [x] You deserve to celebrate!

  Correct

- [x] Andrew sends his heartfelt congratulations to you!

  Correct

- [x] The DeepLearning.AI and Stanford Online teams would like to give you a round of applause!

  Correct